

Deep learning-based super-resolution of climate (forecast) data

Carlos Alberto Gómez Gonzalez
PyconES 2021



**Barcelona
Supercomputing
Center**
Centro Nacional de Supercomputación



A scenic landscape featuring a calm pond in the foreground. In the background, a modern building with a blue, vertically-slatted facade stands behind a dense line of green trees. To the right, a classical-style villa with a brown roof and arched windows is partially visible. The text "About me" is overlaid in white on the center of the image.

About me

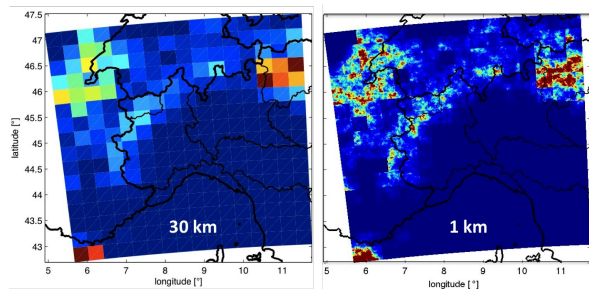
Postdoctoral fellow at the Earth Sciences department of the Barcelona Supercomputing Center (BSC) on Artificial Intelligence (AI) and Machine Learning (ML) projects



About the BSC:


- MareNostrum is the most powerful supercomputer in Spain
- BSC hosts other clusters, such as the Power-CTE, composed of 52 compute nodes each with 4 NVIDIA V100 (Volta) GPUs
- ES department hosts 100+ researchers with expertise in climate science, atmospheric composition among other topics

Climate Science



- Problem definition (need)
- Background and baseline approaches
- Data sources identification
- Validation metrics

Computer vision



- (Re)framing problems
- Cutting edge DL approaches for Earth science applications


AI engineering

```
5 import os
6 import datetime
7 import tensorflow as tf
8 from tensorflow.keras.utils import Progbar
9 import horovod.tensorflow as hvd
10 import numpy as np
11 from dl4ds.resnet_int import resnet_int
12 from dl4ds.resnet_rec import resnet_rec
13 from dl4ds.resnet_spc import resnet_spc
14 from dl4ds.discriminator import residual_discriminator
15 from dl4ds.dataloader import create_pair_hr_lr
16 from dl4ds.utils import (Timing, list_devices,
17                          set_gpu_memory_growth,
18                          set_visible_gpus)
```

- Smart testing and model design/tuning
- Development of robust and efficient code
- Reusability/reproducibility


A scenic view of a pond with a modern building and a traditional building in the background. The modern building has a blue, textured facade, while the traditional building is yellow with a tiled roof. The pond reflects the buildings and the surrounding greenery. Two ducks are swimming in the water.

Introduction



- Spatio-temporal processes
- Common structural prior. Gridded data can be treated as arrays of pixels (images or videos)
- Tasks in CV that relate to a problem in ES:
 - Next frame video prediction → regression
 - Super-resolution → downscaling
 - Object recognition
 - Inpainting → missing data filling
 - Image to image translation → transfer functions, regression

Computer Vision for Earth Science (reality)




Climate scientists

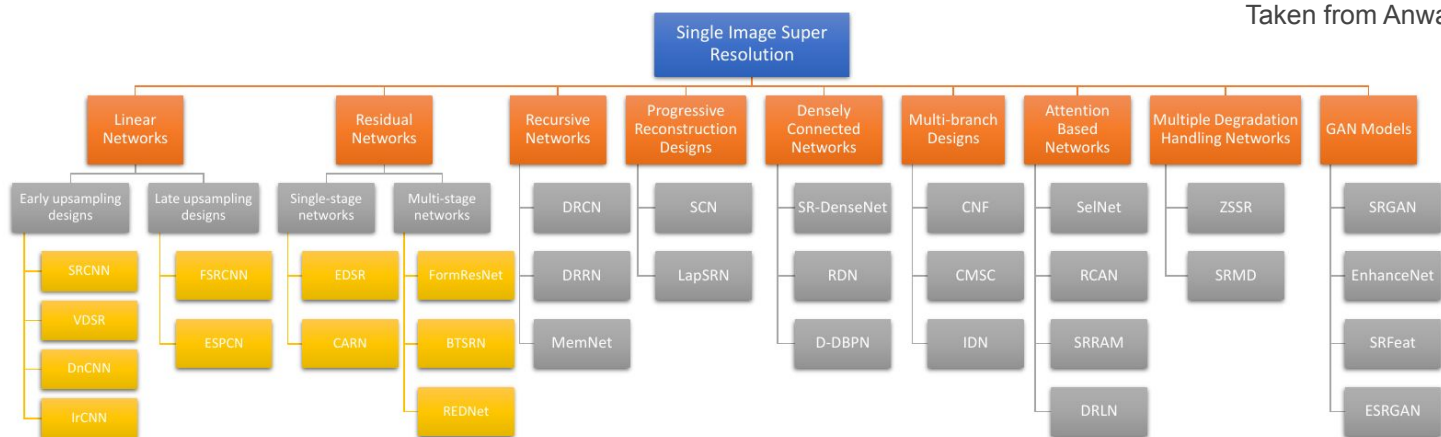
AI/CV/ML specialist

Why super-resolution or downscaling


- Super-resolution aims to convert a given low-resolution image with coarse details to a corresponding high-resolution image with better visual quality and refined details
- Similar goal to that of Statistical Downscaling (SD) techniques in Climate science. SD is much cheaper than running the physical model at higher resolution
- Having more resolution (giving local insights) is important in many applications
- Careful with the terminology:
 - downscaling in climate science == upscaling in CV (from lower to a higher resolution)



- A large number of super-resolution models have been proposed in the field of computer vision



- These ideas have inspired DL-based downscaling methods in climate science, e.g., Vandal et al. 2017, Leinonen et al. 2020, Stengel et al. 2020, Wang et al. 2021, etc



Activation
map (output)

Input

2D convolution using a kernel size of 3 (sliding shadow) with stride of 1 and padding



96 convolutional kernels of size $11 \times 11 \times 3$
learned by the first convolutional layer of
an image classification CNN. From
Krizhevsky et al. 2012

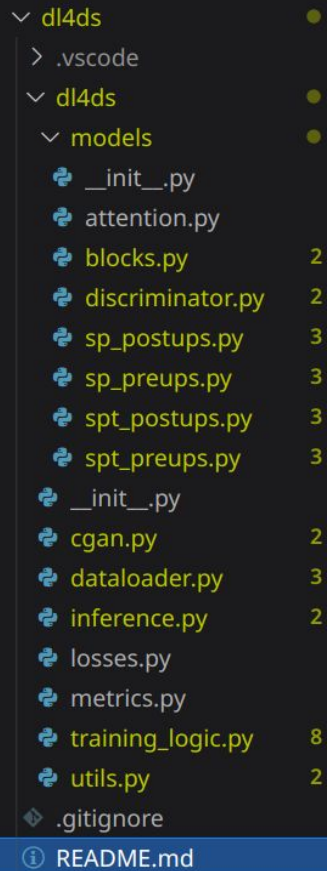
- The convolutional layer is the core building block of a CNN that does most of the computational heavy lifting
- Its parameters consist of a set of learnable filters (image on the top right)
- During the forward pass, we slide each filter across the width and height of the input and compute dot products between the entries of the filter and the input at any position

A scenic view of a pond with a modern building and a classical building in the background. The modern building has a blue, vertically-slatted facade, while the classical building is yellow with a tiled roof. The pond reflects the buildings and the surrounding greenery. Two ducks are visible in the water.

Python and scientific software: DL4DS

- Ugly and tedious job (the most common approach is to publish proofs of concepts)
- (We) scientists suck at programming !!!
 - DRY (Don't Repeat Yourself) **
 - Modularization, abstraction, testing is beneficial in scientific sw development
- Reproducibility crisis is even more relevant in ML/AI
- Python is a vital tool in modern science, data science and ML/AI
- Examples of scientific packages: numpy, scikit-learn, xarray, many other niche packages

** <https://www.earthdatascience.org/courses/earth-analytics/automate-science-workflows/write-efficient-code-for-science-r/>



DL4DS (Deep Learning for DownScaling)

The different design choices can be combined into 48 different network architectures:

Training	Sample type	Backbone block	Upsampling strategy
Supervised (non-adversarial)	Spatial	Plain Conv	Pre-upsampling: interpolation
Adversarial (conditional)	Spatio-temporal	Residual	Post-upsampling: sub-pixel convolution (SPC)
		Dense	Post-upsampling: resize convolution (RC)
			Post-upsampling: deconvolution (DC)


dl4ds > dl4ds > models > blocks.py > ...

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import (Add, Conv2D, ConvLSTM2D,
3                                     SeparableConv2D, BatchNormalization,
4                                     LayerNormalization, Activation,
5                                     SpatialDropout2D, Conv2DTranspose,
6                                     SpatialDropout3D, Concatenate)
7
8 from .attention import ChannelAttention2D
9
10
11 class ConvBlock(tf.keras.layers.Layer):
12     """Convolutional block.
13
14     References
15     -----
16     [1] Effective and Efficient Dropout for Deep Convolutional Neural Networks:
17     https://arxiv.org/abs/1904.03392
18     [2] Rethinking the Usage of Batch Normalization and Dropout:
19     https://arxiv.org/abs/1905.05928
20     """
21     def __init__(self, filters, strides=1, ks_cl1=(3,3), ks_cl2=(3,3),
22                 activation='relu', normalization=None, attention=False,
23                 dropout_rate=0, dropout_variant=None, depthwise_separable=False,
24                 **conv_kwargs):
25         super().__init__()
26         self.normalization = normalization
27         self.attention = attention
28         self.dropout_variant = dropout_variant
29         self.dropout_rate = dropout_rate
30         self.depthwise_separable = depthwise_separable
```


- Written on top of Tensorflow and Keras
- Series of custom layers and model architectures
- Allows easier model design and hyper-parameter optimization
- Distributed training (Horovod)

Daily surface air temperature grids from 1981 to 2020 (~15k time samples) for the Mediterranean region


Observational reference (ERA5 tas 0.25°)



Seasonal forecast (SEAS5 tas 1°, 4x coarser)



Static fields (elevation and land-ocean mask)




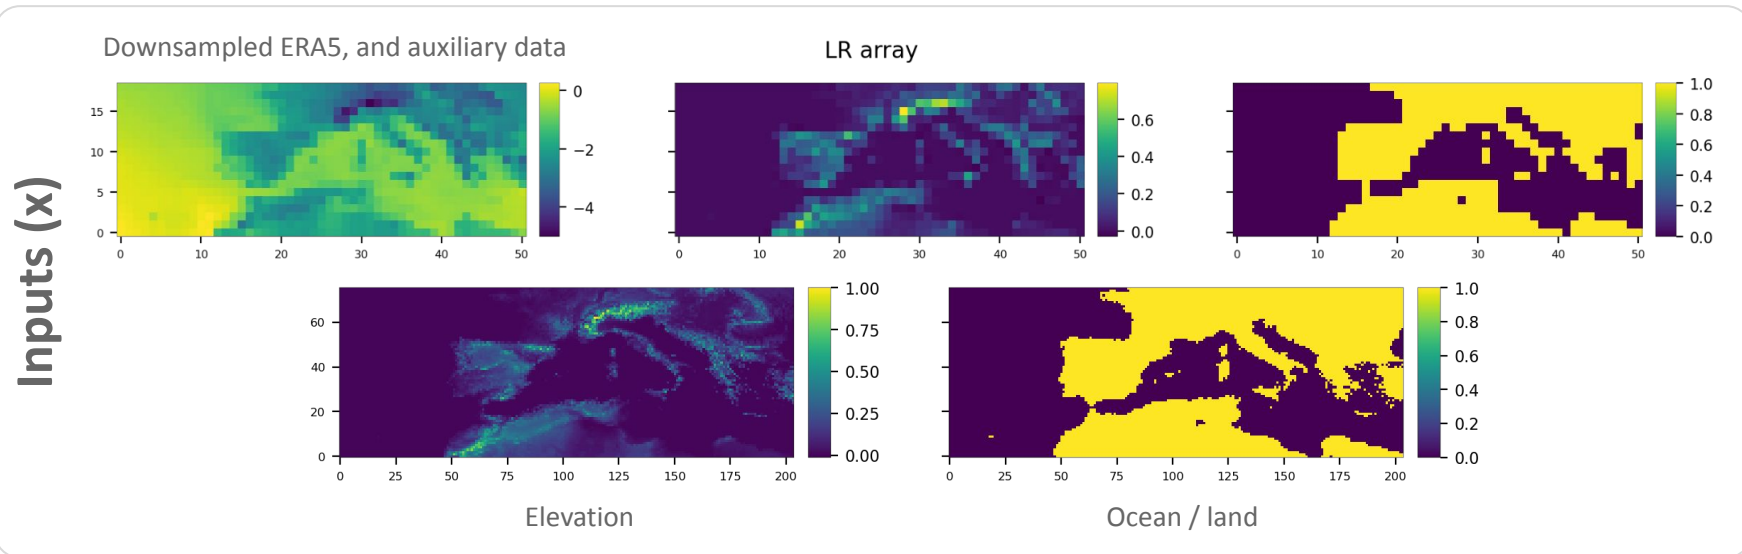
DataGenerator for training samples

```
444 class DataGenerator(tf.keras.utils.Sequence):
445     """
446     A sequence structure guarantees that the network will only train once on
447     each sample per epoch which is not the case with generators.
448     Every Sequence must implement the __getitem__ and the __len__ methods. If
449     you want to modify your dataset between epochs you may implement
450     on_epoch_end. The method __getitem__ should return a complete batch.
451
452     """
453     def __init__(self,
454                 array,
455                 scale=4,
456                 batch_size=32,
457                 patch_size=None,
458                 time_window=None,
459                 topography=None,
460                 landocean=None,
461                 predictors=None,
462                 model='resnet_spc',
463                 interpolation='bicubic',
464                 repeat=None,
465                 ):
466         """
467         Parameters
468         -----
```

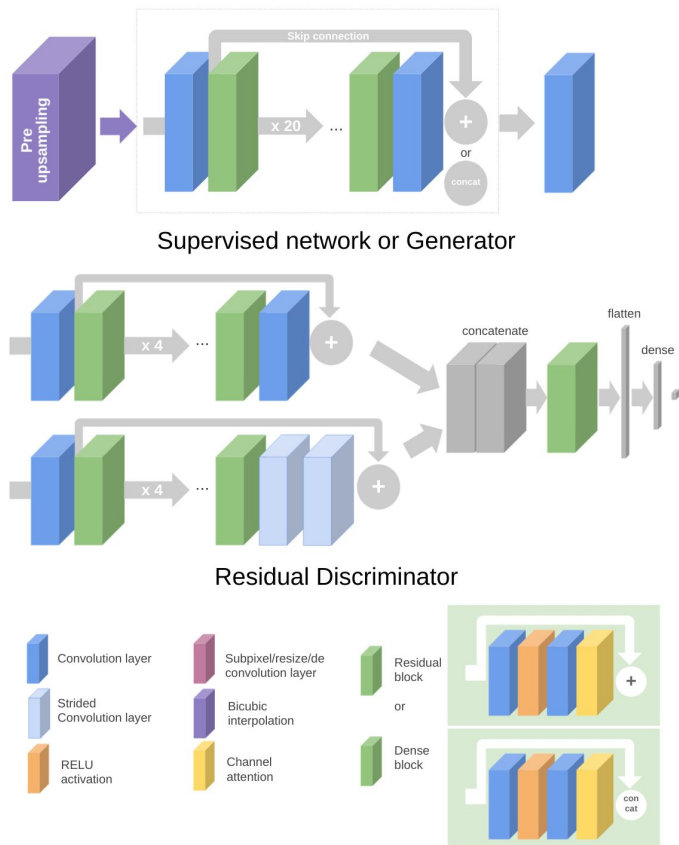
```
517     def __getitem__(self, index):
518         """
519         Generate one batch of data as (X, y) value pairs where X represents the
520         input and y represents the output.
521         """
```

- We subclass keras Sequence
- `__getitem__` returns a batch of samples (X, y) for supervised training
- All the preprocessing is done here (cropping, resizing, slicing, etc)

Training samples (post-upsampling)




Network architectures (pre-upsampling)




- The network can be trained alone in a fully supervised manner
- Both networks, generator and discriminator can be trained in a conditional adversarial fashion
- For a pre-upsampling architecture the images are upsampled via bicubic interpolation before entering the network
- The network feature learning is done in HR and is more computationally expensive


Network architectures (post-upsampling)




Supervised network or Generator




Residual Discriminator



- For a post-upsampling architecture the images are not before entering the network
- The upsampling is performed as a layer of the network itself
- The network feature learning is done in LR and is more computationally efficient
- Deconvolution:



- Subpixel convolution:



Taken from Wang et al., 2020


Taken from Wang et al., 2020


```
GEN_PARAMS = dict(n_filters=32, n_blocks=4, dropout_rate=0.2, dropout_variant='spatial')  
DISC_PARAMS = dict(n_filters=32, n_res_blocks=4)
```

```
model = dds.CGANTrainer(  
    'resnet_spc',  
    data_train,  
    data_test,  
    scale=SCALE,  
    topography=TOPO,  
    landocean=LAOC,  
    predictors_train=None,  
    predictors_test=None,  
    interpolation='bicubic',  
    patch_size=None,  
    batch_size=16,  
    epochs=100,  
    loss='mae',  
    learning_rates=(2e-4, 2e-4),  
    save=True,  
    savecheckpoint_path=None,  
    verbose=1,  
    device='GPU',  
    save_loss_history=False,  
    generator_params=GEN_PARAMS,  
    discriminator_params=DISC_PARAMS)
```

```
model.run()
```

- RELU activations
- Adam optimizer
- Spatial dropout (0.2)
- Several losses are implemented
- Example of a learning curve for a supervised CNN with a MAE loss:





Super resolved temperature

Example from the holdout dataset (samples not used during training)

Groundtruth temperature (20 time steps)

CNN output (from downsampled GT temperature)




A scenic view of a pond with a modern building and a classical building in the background. The modern building has a blue, textured facade, while the classical building is yellow with a red-tiled roof. The pond is surrounded by lush green trees, and two ducks are visible in the water. The text "Application to seasonal forecasts" is overlaid in the center.

Application to seasonal forecasts


Inference on seasonal forecast data

Example from one SEAS5 member at the original low resolution using a given DL4DS network (resnet spc) trained on observational ERA5 temperature data

One SEAS5 (seasonal forecast)



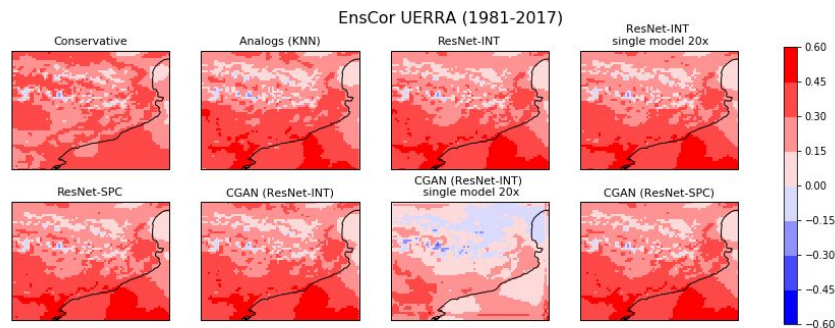

CNN output from the SEAS5 grid on the left




Of course, visual inspection is not enough

Climate science metrics for forecast verification are important for evaluating if the super-resolved seasonal forecast has degraded/improved skill

This is outside of the scope of this talk



- AI/ML has great potential in Earth sciences
- Python rocks!
- Scientific software is key in modern research
- Super-resolution techniques might help improve low resolution seasonal climate forecasts
- DL4DS is part of a publication in preparation and has not yet been open sourced





**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Thanks!

Acknowledgements:

The research leading to these results has received funding from the European Union's Horizon 2020 research and innovation programme under the Marie Skłodowska-Curie grant agreement H2020-MSCA-COFUND-2016-754433 and the EU H2020 Framework Programme under grant agreements n° GA 823988 (ESiWACE-2), GA 869575 (FOCUS-Africa) and GA 869565 (VitiGEOSS).

carlos.gomez@bsc.es